ARAF Insurance Underwriting Crosswalk

Mapping ARAF Dimensional Assessment to D&O, Tech E&O, Cyber, and Professional Indemnity Underwriting

Agentic Risk Architecture Framework, Version 3.0 | Institute for Autonomous Governance Pty Ltd

Disclaimer and Copyright Notice

No Reliance

This document is provided for informational and educational purposes only. It is intended to assist underwriters, brokers, risk officers, and other insurance market participants in understanding how the ARAF dimensional assessment framework may be applied to insurance underwriting workflows. Readers should not rely on this document as a substitute for independent analysis, professional judgment, or advice from appropriately qualified professionals. No representation or warranty, express or implied, is made by the Institute for Autonomous Governance Pty Ltd regarding the accuracy, completeness, currency, or applicability of the information contained herein to any particular organisation, deployment, transaction, or circumstance.

Not Legal, Insurance, or Professional Advice

Nothing in this document constitutes legal, insurance, actuarial, regulatory, or other professional advice. The information is general in nature and has not been prepared with reference to the specific circumstances, risk profile, regulatory environment, or objectives of any reader or organisation. This document should not be relied upon for specific underwriting decisions, coverage determinations, premium calculations, claims assessments, or legal conclusions. Readers are strongly encouraged to seek advice from appropriately qualified legal counsel, licensed insurance and reinsurance professionals, actuaries, and other relevant advisors before making any decision or taking any action based on the content of this document. References to legal authorities, regulatory standards, and market practices are provided for illustrative purposes only and do not constitute legal opinion or analysis.

No Liability

To the fullest extent permitted by applicable law, the Institute for Autonomous Governance Pty Ltd, and its authors, directors, officers, employees, contractors, and affiliates disclaim all liability for any loss, damage, cost, or expense of any kind, whether direct, indirect, consequential, or otherwise, arising out of or in connection with any reliance on this document or any errors, omissions, or inaccuracies in its content. This disclaimer applies regardless of the form of action and whether in contract, tort, negligence, or otherwise.

Copyright

Copyright 2026 Institute for Autonomous Governance Pty Ltd. All rights reserved.

This document and its contents are the intellectual property of the Institute for Autonomous Governance Pty Ltd. No part of this document may be reproduced, stored in a retrieval system, distributed, transmitted, published, or otherwise made available in any form or by any means, whether electronic, mechanical, photocopying, recording, or otherwise, without the prior written permission of the Institute for Autonomous Governance Pty Ltd, except that brief quotations may be included for the purposes of review, commentary, or academic citation, provided that the source is appropriately acknowledged. Requests for permission to reproduce or distribute this document should be directed to the Institute for Autonomous Governance Pty Ltd.

At a Glance

Element	Summary
GBI ≤1.75 (Certified)	Strong governance. Preferred pricing class. Strongest claims defensibility.
GBI 1.75–2.50 (Compliant)	Adequate governance with conditions. Standard pricing. Monitor GCI at renewal.
GBI > 2.50 (Not Compliant)	Inadequate governance. Decline, exclude, surcharge, or condition on remediation.
Active multiplier	Compound governance failure. Enhanced review regardless of composite GBI.
Tier 1 evidence	Infrastructure-generated, contemporaneous. Strongest claims defence.
Tier 3 / absent evidence	Governance record cannot be reconstructed. Weakest defence position.
Key D&O dimensions	D1 (autonomy level), D4 (liability architecture), D5 (exit capacity), D6 (governance drift), GCI, Evidence Tier
Key E&O dimensions	D1 (committal vs. advisory), D2 (data sensitivity), D3 (contract coverage), GCI
Key Cyber dimensions	D2 (data exposure), D3 (DPA adequacy), D5 (vendor dependency), D6 (drift)
Key PI dimensions	D1 (review depth), D4 (liability allocation), Evidence Tier

Purpose

Insurance underwriters pricing coverage for organisations deploying autonomous systems currently lack a standardised governance signal. The Lloyd's Market Association's January 2026 survey confirmed this condition: the survey was conducted because the actuarial data that standard loss modelling requires does not yet exist for AI exposures. Underwriters are pricing autonomous system risk against class averages, general industry assessments, and management representations rather than against the governance posture of specific deployments.

This crosswalk maps ARAF dimensional assessment output to the underwriting workflow for the four insurance lines most exposed to autonomous system deployment: directors and officers liability (D&O), technology errors and omissions (tech E&O), cyber insurance, and professional indemnity. It is designed for the underwriter evaluating a submission, the broker preparing the placement, and the GC or risk officer managing the organisation's insurance programme.

For underwriters: Part 1 (GBI score as underwriting input) and the underwriting signal column in each Part 2 table provide the decision inputs for submission evaluation and pricing. Part 3 maps ARAF outputs to submission, renewal, and claims workflows.

For brokers: Part 4 (conditioning coverage on ARAF certification) provides the sample coverage language, governance-conditional endorsement structure, and exclusion carve-back mechanism for placement. The dimensional mapping in Part 2 supports the submission narrative by translating governance posture into underwriting-relevant signals for each insurance line.

For general counsel and risk officers: the claims defensibility analysis in Part 3 and the evidence tier implications throughout Part 2 support insurance programme design and pre-loss governance positioning.

The crosswalk demonstrates how the GBI score and dimensional profile function as underwriting inputs: enabling risk segmentation by governance posture, informing coverage scope and pricing, and determining claims defensibility before an adverse outcome occurs.

The Underwriting Problem ARAF Addresses

The insurance market's current response to autonomous system risk is exclusion rather than pricing. Standardised policy endorsements issued by the Insurance Services Office (ISO CG 40 47 and CG 40 48) remove AI-generated output from coverage under commercial general liability policies. Individual Lloyd's market carriers have introduced AI-specific exclusions across affected lines. Carriers have introduced absolute AI exclusions across D&O, E&O, and fiduciary liability lines.

Exclusion is the rational market response when risks cannot be segmented. An underwriter who cannot distinguish between a well-governed deployment and a poorly-governed deployment must price both at the worst-case rate or exclude the exposure entirely. Governance classification changes this equation. An underwriter who can assess governance posture through a standardised, independently verified signal can segment risk, price to governance quality, and offer coverage where exclusion was previously the only option.

The GBI provides that signal. The dimensional profile provides the risk segmentation detail. The evidence tier classification provides claims defensibility assessment before the claim arises.

The mechanism for this transition is direct. Where an insurer has applied an AI exclusion (whether ISO CG 40 47/48, a carrier-specific AI exclusion, or equivalent), an ARAF assessment provides the governance signal that enables the insurer to replace the blanket exclusion with a governance-conditional endorsement. The sequence is: (1) AI exclusion applied as the default coverage position; (2) insured presents a current ARAF assessment at Certified or Compliant tier; (3) insurer replaces or narrows the exclusion through a governance-conditional endorsement that covers AI exposure for assessed systems meeting the specified governance threshold. This is how governance classification converts exclusion into pricing.

Part 1: The GBI Score as an Underwriting Input

GBI Score Range	Certification Tier	Underwriting Implication Across All Lines
1.0 to 1.75	ARAF Certified	Strong governance posture. The deployment has been independently assessed as having governance architecture that supports institutional reliance. Underwriting can price to a preferred governance class. Claims defensibility is highest: Tier 1 evidence, operational GCI, documented authority boundaries. Coverage can be offered with confidence that the insured's governance record will support the defence.
1.75 to 2.50	ARAF Compliant	Adequate governance posture with conditions. The deployment has governance architecture in place but with identified areas requiring enhanced monitoring. Standard pricing applies with governance conditions noted. Claims defensibility is adequate but conditional on GCI maintenance. Underwriter should monitor GCI at renewal.
Above 2.50	Not ARAF Compliant	Governance architecture is inadequate for institutional reliance. Remediation is required. Underwriting options: decline coverage for the AI-specific exposure, offer coverage with AI-specific exclusions, offer coverage with a materially higher premium reflecting the governance gap, or offer coverage conditional on achieving Compliant certification within a defined period. Claims defensibility is weak: the governance record may not support the insured's position.
Active multiplier (any)	N/A (compound risk)	The multiplier identifies a compound governance failure requiring specific underwriting attention. Systemic Escalation (D1 + D4): highest D&O exposure. Infrastructure Collapse (D3 + D1): highest E&O exposure from contractual gaps. Leverage Collapse (D5 + D4): highest concentration risk. An active multiplier should trigger enhanced underwriting review regardless of the composite GBI score.

Part 2: Dimensional Mapping to Insurance Lines

D&O (Directors and Officers Liability)

Underwriting question: Can the board demonstrate it governed the autonomous systems producing consequential decisions?

The accountability standard for director governance of autonomous systems is established across major common law jurisdictions. In Australia, the institutional accountability test—established across ASIC v Bekier [2026] FCA 196, ASIC v Healey [2011] FCA 717, and ASIC v Cassimatis [2016] FCA 1023—holds that officers with operational responsibility must demonstrate substantive engagement and action, not merely governance documentation. Non-executive directors may rely on management absent information that would awaken suspicion; the duty requires that governance information capable of triggering that suspicion exists and reaches the board. In the United States, the Caremark/Marchand line imposes equivalent monitoring obligations for mission-critical risks. Comparable director accountability standards apply under the UK Companies Act 2006 (ss 172–174) and equivalent statutes across Commonwealth and EU jurisdictions. These authorities establish the legal standard against which D&O claims arising from autonomous system governance failures will be assessed.

ARAF Dimension	Relevance to Underwriting	What the Underwriter Should Examine	Underwriting Signal
D1: Autonomy Gradient	The higher the autonomy level, the greater the board's exposure if it cannot demonstrate adequate oversight. A system operating at high autonomy (D1 ≥4) with human oversight that is supervisory rather than authorial creates the governance failure the institutional accountability test identifies: governance structures that exist but are not substantively exercised.	D1 score and the oversight adequacy sub-assessment. A D1 score of 1.0–2.0 indicates meaningful human oversight with documented authority boundaries. A score above 3.0 indicates oversight gaps. A score of 4.0 or above indicates the system operates with authority levels that exceed the governance architecture's capacity to oversee it.	D1 above 3.0: elevated D&O exposure. The board may be unable to demonstrate it understood and governed the authority delegation. D1 at or below 2.0: governance architecture supports the board's position that oversight was adequate.
D4: Liability Architecture	D4 determines whether the organisation has defined what happens when autonomous action produces adverse consequences. A board that deployed a system without a liability architecture for the decisions it produces cannot answer "what did they do" under the institutional accountability test (established across Bekier, Healey, and Cassimatis; the three-question test: who held authority, what did they know, what did they do).	D4 score and the Autonomous Event Level 3 (AE3) classification, which identifies decisions where the autonomous system commits the organisation to actions no human approved. A D4 score of 1.0–2.0 indicates defined liability boundaries, documented caps, and insurance coverage aligned to the deployment scope. A score above 3.0 indicates material gaps. A score of 4.0 or above indicates no formal liability architecture exists.	D4 above 3.0: material D&O exposure. The board authorised a deployment without a liability architecture to absorb its consequences. D4 at or below 2.0: liability architecture is defined and documented.
D1 + D4 combined (Systemic Escalation multiplier)	The Systemic Escalation multiplier (D1 ≥4 AND D4 ≥4) identifies the highest-risk governance profile for D&O purposes: high-speed autonomous decisions producing consequences that no liability framework is designed to absorb. This is the compound condition where D&O claims are most likely to succeed because the board cannot demonstrate either adequate oversight (D1) or adequate liability preparation (D4). This is the governance equivalent of issuing a financial guarantee without a capital reserve.	Whether the Systemic Escalation multiplier is active. If active, the underwriter is looking at a deployment where the board delegated high autonomous authority without a corresponding liability architecture.	Systemic Escalation active: the highest D&O risk classification. The board's governance position is weakest precisely where the consequences are greatest.
D5: Commercial Leverage	D5 identifies whether the organisation can exit or substitute the autonomous system if governance remediation is required. A board that authorised a deployment it cannot exit has delegated authority without retained governance capacity. For	D5 score and remediation resistance sub-assessment. A D5 score above 3.0 indicates structural dependency on an AI vendor or system the organisation cannot practically exit. Combined with D4 ≥3 (Leverage Collapse	D5 above 3.0: the board's ability to remediate governance failures is commercially constrained. D&O exposure increases because the defence that "the board took corrective action" may not be available if

	D&O purposes, this is a governance judgment finding: the board's remediation options are constrained by commercial dependency.	multiplier), the organisation is both locked in and unprotected.	corrective action requires exiting a system the organisation cannot leave.
D6: Adaptive Stability	D6 determines whether the governance posture assessed at policy inception will persist over the policy period. A system assessed as well-governed at binding may drift through model updates, scope expansion, or authority threshold changes. The D&O risk profile at binding may not be the D&O risk profile at claim.	D6 score and the Drift Detection component. An underwriter should assess whether the governance architecture has mechanisms to detect and report changes that alter the D&O risk profile during the policy period.	D6 above 3.0: governance may not maintain currency over the policy period. Consider more frequent reporting obligations or a mid-term review condition. D6 at or below 2.0: governance architecture includes monitoring and reassessment mechanisms that maintain currency.
GCI (Governance Coherence)	The GCI determines whether governance was exercised during the policy period, not only whether governance structures existed. A D&O claim will turn on whether the board can produce evidence of governance action, not merely governance documentation.	GCI composite and the Control Exercise component specifically. A strong GBI with a weak GCI indicates governance on paper that was not exercised in practice—the Bekier failure condition. An underwriter should treat a GCI failure as equivalent to a control environment failure in traditional D&O analysis.	GCI below conformance threshold: the board has governance documentation but cannot demonstrate governance action. This weakens the insured's defence position and increases the underwriter's expected loss severity.
Evidence Tier (EIS-01)	The evidence tier classification determines claims defensibility for D&O. A board that can produce Tier 1 evidence (infrastructure-generated, contemporaneous governance records) has the strongest defence position. A board that cannot reconstruct the governance conditions at the time of the autonomous action that produced the loss has a materially weaker defence.	Evidence tier classification across D1 and D4 primarily. Tier 1 evidence supports the insured's position that governance was exercised. Tier 3 or absent evidence means the governance conditions cannot be reconstructed, weakening the insured's position under Bekier and potentially constituting a failure to mitigate.	Evidence at Tier 1: D&O claims defensibility is supported by contemporaneous governance records. Evidence at Tier 3 or absent: the board cannot demonstrate governance action at the time of the event. Expected loss severity increases.

Tech E&O (Technology Errors and Omissions)

Underwriting question: Will the AI system's outputs produce professional liability exposure?

ARAF Dimension	Relevance to Underwriting	What the Underwriter Should Examine	Underwriting Signal
D1: Autonomy Gradient	The autonomy level determines whether the system produces advisory outputs (lower E&O exposure, human acts on recommendation) or committal outputs (higher E&O exposure, system executes action directly). A system that commits transactions, deploys code, or sends communications autonomously creates E&O exposure at the execution layer.	D1 score and the commitment authority sub-assessment. Committal authority (the system can execute consequential actions without human approval) creates a different E&O exposure profile from advisory authority (the system recommends, a human decides).	D1 indicating committal authority: E&O exposure attaches at the execution layer. D1 indicating advisory authority: E&O exposure attaches at the recommendation layer. The premium and coverage structure should reflect this distinction.
D2: Data Sensitivity Exposure	The data categories the system processes determine the magnitude of E&O exposure if the system produces erroneous outputs. A system processing financial data, health records, or legal analysis produces higher-consequence errors than a system processing internal scheduling data.	D2 score, distinguishing between operational data sensitivity (what the system processes in production) and training data provenance (what the system was trained on). An E&O claim may arise from either: erroneous outputs caused by operational data quality issues, or systematic bias caused by training data deficiencies.	D2 at or below 2.0: data sensitivity is lower, and E&O exposure from data-driven errors is bounded.
D3: Contract Infrastructure	D3 determines whether the contract between the AI vendor and the	D3 score and the liability provisions sub-assessment. Does	D3 above 3.0 combined with D1 at or above 3.0

	deploying organisation allocates E&O liability for AI-generated outputs, or whether it defaults to standard limitation clauses not designed for autonomous system errors.	the vendor's contract explicitly address liability for errors in AI-generated outputs, or does it limit liability to standard service availability metrics? A gap between what the system can do (D1) and what the contract covers (D3) is a direct E&O underwriting finding.	(Infrastructure Collapse multiplier: D3 ≥4 AND D1 ≥3): contractual liability allocation does not cover the autonomous action scope. E&O exposure may fall entirely on the deploying organisation.
GCI (Governance Coherence)	The GCI determines whether the governance controls that supported the E&O risk assessment at binding were still operating at the time of the alleged error. An E&O claim arising from AI-generated output will turn on whether the deployer can demonstrate that governance controls were operating when the erroneous output was produced.	GCI composite and Control Exercise component. A strong GBI with a weak GCI indicates the deployer had governance architecture in place but was not exercising it when the error occurred. For E&O purposes, this is the difference between a system error (potentially defensible) and a governance failure (harder to defend).	GCI below conformance threshold: governance controls were not operating at the required cadence. E&O claims defensibility is weakened because the deployer cannot demonstrate the error occurred despite active governance.

Cyber Insurance

Underwriting question: What is the data exposure profile of this AI deployment?

ARAF Dimension	Relevance to Underwriting	What the Underwriter Should Examine	Underwriting Signal
D2: Data Sensitivity Exposure	D2 directly maps to the data exposure assessment cyber underwriters perform for any technology deployment. The additional complexity for AI systems: the data exposure surface includes both operational data (what the system processes in production) and training data (what the system has ingested and potentially memorised).	D2 score and the dual classification (operational vs training data). A cyber underwriter should assess both exposure surfaces. An AI system that processes personal data operationally AND was trained on proprietary data creates a compound data exposure that standard cyber policies may not be priced for.	D2 above 3.0: elevated data exposure requiring AI-specific pricing consideration. D2 at or below 2.0: data exposure within standard cyber risk parameters.
D3: Contract Infrastructure	D3 identifies whether the data processing agreements across the AI deployment chain address the data categories the system actually processes. A gap between D2 (data exposure) and D3 (contractual coverage) is a cyber underwriting finding.	D3 score with specific attention to data processing agreement adequacy. Does the DPA cover the data categories identified in D2? Does it address AI-specific processing (inference, embeddings, retrieval-augmented generation) or only standard data storage and transmission?	Gap between D2 and D3: the organisation is processing data through AI systems under contracts not designed for that processing. This is the cyber equivalent of silent AI risk: AI exposure that exists within the policy portfolio but is neither explicitly covered nor explicitly excluded, paralleling the "silent cyber" condition the market addressed through LMA clauses in 2019–2020.
D5: Commercial Leverage	D5 identifies whether the organisation's incident response capability is constrained by vendor dependency. An organisation locked into an AI vendor it cannot exit may be unable to isolate, contain, or remediate an AI-specific data exposure incident without vendor cooperation.	D5 score combined with D3 (contractual incident response obligations). Does the vendor's contractual obligations include AI-specific incident response cooperation, or does the organisation depend on vendor goodwill to manage a cyber event arising from the AI deployment?	D5 above 3.0 combined with D3 above 3.0: the organisation's cyber incident response for AI-related events depends on a vendor it has limited contractual leverage over. Incident containment timelines may be extended by vendor dependency.
D6: Adaptive Stability	D6 is relevant to cyber underwriting because model updates, data pipeline changes, and scope expansion can alter the system's data exposure profile between assessment periods. A system assessed as low-risk at deployment may become high-risk after a model	D6 score and the Drift Detection component. Does the governance architecture maintain current visibility into the system's data exposure as the system evolves?	D6 above 3.0: the governance architecture may not detect changes to the data exposure profile. Cyber pricing should account for the possibility that the risk profile at binding is not the risk profile at claim.

	update that introduces new data categories into the inference pipeline.		

Professional Indemnity

Underwriting question: Is the professional relying on AI-generated analysis for client-facing advice?

ARAF Dimension	Relevance to Underwriting	What the Underwriter Should Examine	Underwriting Signal
D1: Autonomy Gradient	Determines whether the AI system is producing analysis the professional relies upon (advisory) or producing client-facing outputs directly (committal). A professional who delivers AI-generated analysis to a client without substantive review is in the "supervisory not authorial" condition: the system produced the substantive work, the professional approved its delivery.	D1 score. For PI purposes, even advisory-level autonomy creates exposure if the professional cannot demonstrate they substantively reviewed the AI-generated analysis rather than merely approving its output. The standard established in ASIC v Healey [2011] FCA 717 applies: engagement must be substantive, not merely approval of receipt.	D1 above 2.0 in a professional services context: the professional may be delivering AI-generated work product under conditions where their review is supervisory rather than authorial. PI exposure increases because the professional's defence depends on demonstrating they exercised independent judgment.
D4: Liability Architecture	D4 determines whether liability for errors in AI-generated professional advice is allocated between the AI vendor, the professional firm, and the client, or whether it defaults entirely to the professional through standard professional liability.	D4 score and the liability allocation across the advisory chain. Has the professional firm allocated AI-specific liability to the vendor, or absorbed it within its own PI coverage?	D4 above 3.0 in a professional services context: the firm is delivering AI-assisted professional advice without a liability architecture designed for AI-generated errors. PI claims arising from AI-generated analysis will be absorbed entirely by the firm's PI programme.
Evidence Tier (EIS-01)	The evidence tier classification is directly relevant to PI claims defensibility. A professional who can produce Tier 1 evidence records of the AI system's analysis, the data it relied upon, and the governance controls that were active has a defensible position. A professional who cannot reconstruct the AI system's analytical process does not.	Evidence tier classification for the relevant dimensions (D1 and D2 primarily). Tier 1 evidence supports the professional's defence that the AI-assisted work was produced under governed conditions. Tier 3 or absent evidence means the professional cannot demonstrate what the AI system did or why.	Evidence at Tier 1: the professional's PI defence is supported by contemporaneous governance records. Evidence at Tier 3 or absent: the professional cannot reconstruct the analytical basis of the advice they delivered. PI claims defensibility is materially weakened.

Part 3: Using the Crosswalk in the Underwriting Workflow

At submission / new business

The broker presents the ARAF assessment report as part of the submission, alongside SOC 2 attestation and other governance documentation. The underwriter reviews the GBI score, dimensional profile, multiplier status, and evidence tier classification. The dimensional profile allows the underwriter to assess risk by governance dimension rather than by class average. Two submissions with identical revenue, headcount, and AI deployment descriptions but different GBI scores represent materially different risks. The GBI score is the signal that enables that segmentation.

At renewal

The underwriter compares the current GBI score against the prior period. Has the governance posture improved, declined, or remained stable? Has the GCI been maintained (governance is still operating) or has it degraded (governance has lapsed)? Have any reassessment triggers fired since the last assessment? Has the dimensional profile shifted in ways that alter the risk for specific insurance lines (e.g., D1 increased due to scope expansion, altering the D&O exposure)? The GBI trend over time is as informative as the point-in-time score: an organisation with an improving GBI trajectory represents a different risk from one with a stable or declining trajectory.

At claim

The evidence tier classification determines claims defensibility. The claims handler examines: what governance records existed at the time of the incident? At what evidence tier were they classified? Does the governance record support or undermine the insured's position that the deployment was governed at the time the adverse outcome occurred?

A Tier 1 governance record (infrastructure-generated, contemporaneous, not retrospectively assembled) supports the insured's defence. A Tier 3 or absent record means the governance conditions cannot be reconstructed, which weakens the insured's position under the Bekier standard and may constitute a failure to mitigate that affects coverage. The interaction with the insured's duty of disclosure (Insurance Contracts Act 1984 (Cth) s 21; Insurance Act 2015 (UK) s 3, and equivalent obligations in other jurisdictions) should also be considered: where ARAF status is material to the risk and the insured's governance posture has deteriorated since inception, the failure to disclose that deterioration may independently affect the insurer's position.

The evidence tier report, produced at the time of ARAF assessment, is the artefact that determines this outcome. It exists before the claim. That is its value.

Part 4: Conditioning Coverage on ARAF Certification

The crosswalk supports the first stage of the propagation mechanism described in Decision Infrastructure Governance: The Institutional Case for Trust Architecture: insurers condition coverage on governance certification. A coverage condition could be structured as:

> *"The Insured shall maintain, for each autonomous decision system producing or materially contributing to consequential organisational decisions, an ARAF assessment with a GBI score at or below 2.50 (ARAF Compliant) conducted by an assessor accredited by the Institute for Autonomous Governance. The assessment report, including dimensional profile, GCI status, and evidence tier classification, shall be provided to the Insurer at inception and at each renewal. Material changes to the GBI score or certification status shall be notified to the Insurer within the 60-day notification period defined in the reassessment trigger requirements of the current ARAF standard. Where a multiplier condition (Systemic Escalation, Infrastructure Collapse, or Leverage Collapse) activates during the policy period, the Insured shall notify the Insurer within 15 business days. The Insurer may, upon receipt of such notification, require supplementary assessment and may adjust terms by endorsement. Failure to maintain ARAF Compliant status for a system previously assessed shall entitle the Insurer to amend coverage terms, adjust premium, or impose additional conditions at the next renewal."*

This language: specifies the governance standard (ARAF Compliant); defines the scope (consequential decision systems, not all AI); requires independent assessment (accredited assessor); establishes ongoing obligations (renewal, notification); creates defined mid-term notification triggers for compound governance failures; and creates a consequence for governance degradation (coverage terms affected) without creating an automatic coverage exclusion that would defeat the purpose of governance-linked underwriting.

The coverage condition converts the governance signal from a risk assessment input into a policy term. That conversion is the mechanism through which governance standards become commercially mandatory: not because a regulator requires them, but because the insurer does.

Governance-Conditional Endorsement (AI Exclusion Carve-Back)

Where the policy includes an AI exclusion clause (whether ISO CG 40 47/48, a carrier-specific AI exclusion, or equivalent), the following endorsement structure enables the insurer to restore coverage for governed deployments:

> *"Notwithstanding [AI Exclusion Clause Reference], coverage shall extend to loss arising from or in connection with the operation of an autonomous decision system where, at the time of the event giving rise to the loss: (a) the system was the subject of a current ARAF assessment with a GBI score at or below [1.75 for Certified / 2.50 for Compliant] conducted by an assessor accredited by the Institute for Autonomous Governance; (b) the system's Governance Coherence Index was at or above the operational conformance threshold at the most recent assessment; (c) no multiplier condition (Systemic Escalation, Infrastructure Collapse, or Leverage Collapse) was active for the system at the time of the event; and (d) the Insured had complied with the ARAF certification maintenance and notification obligations specified in [coverage condition clause reference]. For the avoidance of doubt, coverage under this endorsement applies only to autonomous decision systems meeting conditions (a) through (d) above. The AI exclusion remains in force for all other AI systems and deployments."*

This endorsement structure converts ARAF certification from a policy condition into a coverage trigger. The insurer retains the AI exclusion as the default position, protecting against ungoverned exposure, while enabling coverage for deployments that meet a defined governance standard. The endorsement creates the commercial incentive the framework describes: organisations that invest in governance architecture gain access to coverage that ungoverned competitors cannot obtain.

Reinsurance Considerations

For placements involving reinsurance (including London market excess-of-loss and proportional treaties), the ARAF assessment report provides the cedant with a standardised governance signal to communicate to reinsurers. A cedant that can demonstrate governance-based risk segmentation across its AI-exposed portfolio (using GBI scores and dimensional profiles for ceded risks) provides reinsurers with the underwriting information needed to price AI exposure based on governance quality rather than portfolio-wide exclusion. As loss data linking GBI scores to claims experience accumulates, insurers may also be able to reflect governance-based risk segmentation in internal capital models, reducing the capital charge allocated to governed AI portfolios.

Martin, Carly. Agentic Risk Architecture Framework (ARAF), Version 3.0. Institute for Autonomous Governance Pty Ltd, 2026.